



JACOBS
UNIVERSITY

The Relational Model

Ramakrishnan & Gehrke, Chapter 3

Relational Database: Definitions

- Technically: **Relation** made up of 2 parts:

- **Schema**: specifies name of relation, plus name and type of each column
 - Ex: Students(sid: string, name: string, login: string, gpa: real)
- **Instance**: a **table**, with rows and columns
 - # rows = **cardinality**, # fields = **degree** / **arity**

does not
change often

changes all
the time

- Mathematically:

- Let A_1, \dots, A_n ($n > 0$) be value sets, called **attribute domains**
- relation $R \subseteq A_1 \times \dots \times A_n = \{ (a_1, \dots, a_n) \mid a_1 \in A_1, \dots, a_n \in A_n \}$

- Can think of a relation as a **set** of rows or **tuples**

- i.e., all rows are **distinct** = no duplicates (hmm...)
- **atomic attribute types** only – no fancies like sets, trees, ...

- **Relational database**: a set of **relations**

Students	sid	name	login	gpa

tuple

attribute

Example Instance of Students Relation



JACOBS
UNIVERSITY

Sid	Name	Login	Gpa
53666	Jones	jones@cs	3.4
53688	Smith	smith@eecs	3.2
53650	Smith	smith@math	3.8

- Cardinality = 3, degree = 4, all rows distinct
- Do all *columns* in a relation instance have to be distinct?

Querying Relational Databases

- A major strength of the relational model: **simple, powerful *querying*** of data
 - Data organised in tables, query results are tables as well
 - Small set of generic operations, work on any table structure
- Query describes **structure of result** ("what"), not algorithm how this result is achieved ("how")
 - data independence, optimizability
- Queries can be written intuitively, and the **DBMS is responsible for efficient evaluation**
 - The key: **precise (mathematical) semantics** for relational queries
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change

SQL, Structured English Query Language



JACOBS
UNIVERSITY

- "find all students with GPA less than 3.6"
 - SELECT *
FROM Students S
WHERE S.gpa < 3.6
- To find just names and logins, replace the first line:
 - SELECT S.name, S.login

sid	name	login	gpa

53666	Jones	jones@cs	3.4
53688	Smith	smith@eecs	3.2
53650	Smith	smith@math	3.8

sid	name	login	gpa

53666	Jones	jones@cs	3.4
53688	Smith	smith@eecs	3.2

name	login

Jones	jones@cs
Smith	smith@eecs

SQL Joins: Querying Multiple Relations



JACOBS
UNIVERSITY

- What does the following query compute?

- ```
SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.sid=E.sid AND E.grade="A"
```

- Given the following instances of Students and Enrolled:

| sid   | name  | login      | gpa |
|-------|-------|------------|-----|
| 53666 | Jones | jones@cs   | 3.4 |
| 53688 | Smith | smith@eecs | 3.2 |
| 53650 | Smith | smith@math | 3.8 |

| sid   | cid         | grade |
|-------|-------------|-------|
| 53831 | Carnatic101 | C     |
| 53831 | Reggae203   | B     |
| 53666 | Topology112 | A     |
| 53688 | History105  | B     |

- we get:

| S.name | E.cid       |
|--------|-------------|
| Jones  | Topology112 |

# DML: Adding and Deleting Tuples

- **insert** a single tuple:
  - INSERT INTO Students( sid, name, login, gpa )  
VALUES ( 53688, 'Smith', 'smith@ee', 3.2 )
- **delete** all tuples satisfying some condition:
  - DELETE FROM Students S  
WHERE S.name = 'Smith'
- **change** all tuples satisfying some condition:
  - UPDATE Students S  
SET gpa = 3.0  
WHERE S.name = 'Smith'

- **Integrity constraint = IC**  
= condition that must be true for any instance of the database
  - e.g., domain constraints
  - ICs are **specified** when schema is **defined**
  - ICs are **checked** when relations are **modified**
- A **legal instance** of a relation is one that satisfies all specified ICs
  - DBMS should not allow illegal instances
- If the DBMS checks ICs, stored data is more faithful to real-world meaning
  - Avoids data entry errors, too!



# Primary Key Constraints

- A set of fields is a **key** for a relation if :
  - 1. No two distinct tuples can have same values in all key fields, and
  - 2. This is not true for any subset of the key.
- Part 2 false → **superkey**
  - If >1 key for relation,  
one of the keys is chosen (by DBA) to be **primary key**
- Example:
  - sid key for Students (what about name?)
  - The set {sid, gpa} is a superkey

# Primary and Candidate Keys in SQL

- Possibly many **candidate keys** (specified using **UNIQUE**), one of which is chosen as the primary key
- *“For a given student and course, there is a single grade”*  
vs.  
*“Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”*
  - Used carelessly, an IC can prevent the storage of database instances that arise in practice!

```
CREATE TABLE Enrolled
(sid CHAR(20)
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid,cid))
```

```
CREATE TABLE Enrolled
(sid CHAR(20)
 cid CHAR(20),
 grade CHAR(2),
 PRIMARY KEY (sid),
 UNIQUE (cid, grade))
```

# Foreign Keys, Referential Integrity

- *Foreign key* = set of fields in one relation that is used to `refer' to a tuple in another relation
  - Must correspond to primary key of the second relation, like a `logical pointer'
- Example: sid is a foreign key referring to Students:
  - Enrolled(sid: string, cid: string, grade: string)
  - If all foreign key constraints are enforced, *referential integrity* is achieved, i.e., no dangling references.
  - Can you name a data model w/o referential integrity?

# Enforcing Referential Integrity

- Consider Students and Enrolled;  
*sid* in Enrolled is a foreign key that references Students
- What should be done if an Enrolled tuple with a non-existent student id is inserted?
  - Reject it
- What should be done if a Students tuple is deleted?
  - Also delete all Enrolled tuples that refer to it
  - Disallow deletion of a Students tuple that is referred to
  - Set *sid* in Enrolled tuples that refer to it to a *default sid*
  - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a *special value NULL*, denoting *'unknown'* or *'inapplicable'*)
- Similar if primary key of Students tuple is updated

# Referential Integrity in SQL

- SQL/92 and SQL:1999 support all 4 options on deletes and updates:
  - Default is **NO ACTION**  
(delete/update is rejected)
  - **CASCADE**  
(also delete all tuples that refer to deleted tuple)
  - **SET NULL / SET DEFAULT**  
(sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT)
```


treat corresponding Enrolled tuple  
when Students (!) tuple is deleted

# Where do ICs Come From?

- based upon the semantics of the real-world enterprise that is being described in the database relations
- can check a database instance to see if an IC is violated, but can **NEVER** infer that an IC is true by looking at an instance
  - An IC is a statement about all possible instances!
  - From example, we know name is not a key, but the assertion that sid is a key is given to us
- Key and foreign key ICs are the most common; more general ICs supported too

# Logical DB Design: ER to Relational

- Entity sets to tables:
  - ER attribute → table attribute  
(can do that because ER constrained to simple types, same as in relational model)
  - Declare key attribute “Primary key”
- Best practice (not followed by book):  
Add “abstract” identifying key attribute
  - No further semantics
  - System generated
  - use only this as primary key & for referencing



```
CREATE TABLE Employees
(ssn CHAR(11),
 name CHAR(20),
 lot INTEGER,
 PRIMARY KEY (ssn))
```

```
CREATE TABLE Employees
(sid INTEGER,
 ssn CHAR(11) UNIQUE,
 ...,
 PRIMARY KEY (sid))
```

# Relationship Sets to Tables



- In translating a relationship set to a relation, attributes of the relation must include:
  - Keys for each participating entity set (as **foreign keys**)
    - *a superkey for the relation*
  - All descriptive attributes

```
CREATE TABLE Works_In
(ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (ssn, did),
 FOREIGN KEY (ssn)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
```




# Review: Key Constraints


- Each dept has at most one manager, according to the key constraint on Manages




1-to-1



1-to-Many



Many-to-1




Many-to-Many

*Translation to  
relational model?  
...see next!*

# ER Diagrams with Key Constraints

- Map relationship to table:

- **did** key now!
- **Separate** tables for Employees and Departments



```
CREATE TABLE Manages
(ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn) REFERENCES Employees,
 FOREIGN KEY (did) REFERENCES Departments)
```


- each department has unique manager  
→ can **combine**  
Manages and Departments

```
CREATE TABLE Dept_Mgr
(did INTEGER,
 dname CHAR(20),
 budget REAL,
 ssn CHAR(11),
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn) REFERENCES Employees)
```

# Participation Constraints in SQL

## ■ Review: Participation Constraints

- Does every department have a manager?  
→ **participation constraint**
- Every did value in Departments table must appear in a row of the Manages table (with non-null ssn value!)




- can capture participation constraints involving one entity set in a binary relationship
  - but little else (w/o CHECK constraints)

```
CREATE TABLE Dept_Mgr
(did INTEGER,
 dname CHAR(20),
 budget REAL,
 ssn CHAR(11) NOT NULL,
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn)
 REFERENCES Employees
 ON DELETE NO ACTION)
```

# Translating Weak Entity Sets

- Review: **weak entity**:  
identifiable uniquely only by *owner* entity
  - one-to-many relationship set  
(1 owner, many weak entities)
  - Weak entity:  
total participation in **identifying** relationship set

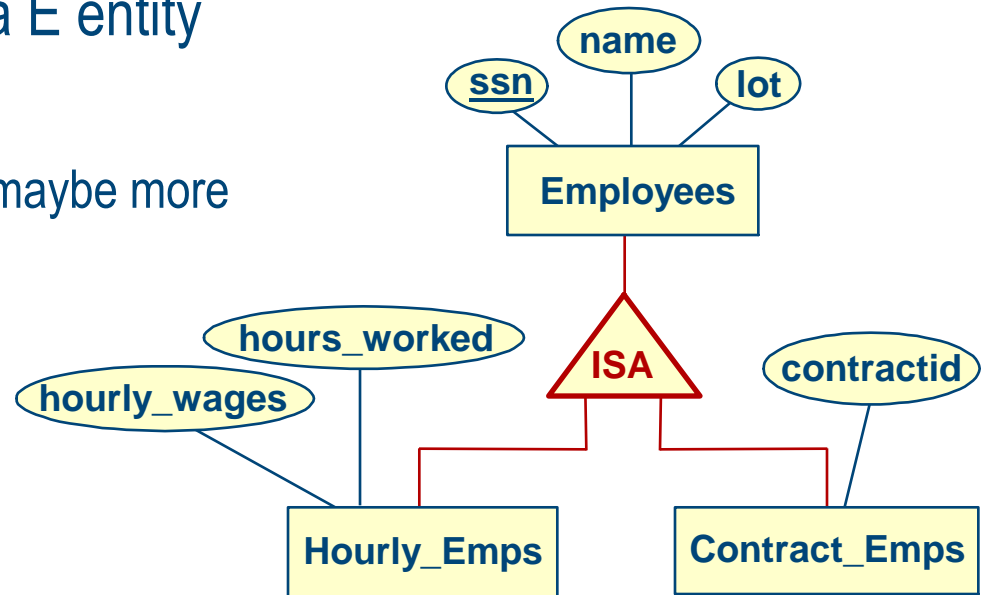


- Weak entity set & identifying relationship set  
→ **single table**
- When owner entity is deleted:  
**delete all owned weak entities**

```
CREATE TABLE Dep_Policy
(pname CHAR(20),
 age INTEGER,
 cost REAL,
 ssn CHAR(11) NOT NULL,
 PRIMARY KEY (pname, ssn),
 FOREIGN KEY (ssn)
 REFERENCES Employees
 ON DELETE CASCADE)
```

# Review: ISA Hierarchies

- **H ISA E**: every H entity is also a E entity  
("H inherits from E")
  - H attributes = E attributes + plus maybe more
  - H **subclass**, E **superclass**
- Mapping to Relations
  - Several choices
  - Constraints determine



# Translating ISA Hierarchies to Relations



JACOBS  
UNIVERSITY

- General approach: **separate relation per entity set**  
→ 3 relations: Employees, Hourly\_Emps, Contract\_Emps
  - Every employee recorded in Employees
  - For hourly emps, extra info recorded in Hourly\_Emps (hourly\_wages, hours\_worked, ssn)
  - **must delete** Hourly\_Emps tuple if referenced Employees tuple is deleted
  - Queries involving all employees easy,  
those involving just Hourly\_Emps require a **join** to get some attributes
- Alternative: **relations only entity sets with instances**  
→ 2 relations: Hourly\_Emps, Contract\_Emps
  - Hourly\_Emps: ssn, name, lot, hourly\_wages, hours\_worked
  - Each employee must be in one of these two subclasses

- **view** is just a relation, but we **store definition** rather than a set of tuples
  - **CREATE VIEW** YoungActiveStudents (name, grade)  
**AS** SELECT S.name, E.grade  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid and S.age < 21
- Views can be dropped using **DROP VIEW**
  - **DROP TABLE** if there's a view on the table? → options

# Views and Security

- Views useful for personalized information (or a summary), while **hiding details** in underlying relation(s)
  - Given YoungStudents, but not Students or Enrolled, we can find students who are enrolled
  - ...but not the cid's of the courses they are enrolled in